

Task and Motion Planning using Learning from Demonstration and Reinforcement Learning

Dominik Urbaniak

Master Thesis – Final Presentation

Dr. Alejandro Agostini

Human-centered Assistive Robotics

Technische Universität München



Dirty Dishes

Washing dishes is one common task in our daily life

→ A dishwasher is helpful, but still requires to be loaded

Requirements:

Carefully picking and placing objects of various size and shape in a meaningful order.

→ Task and motion planning problem



<https://home.howstuffworks.com/dishwasher.htm>

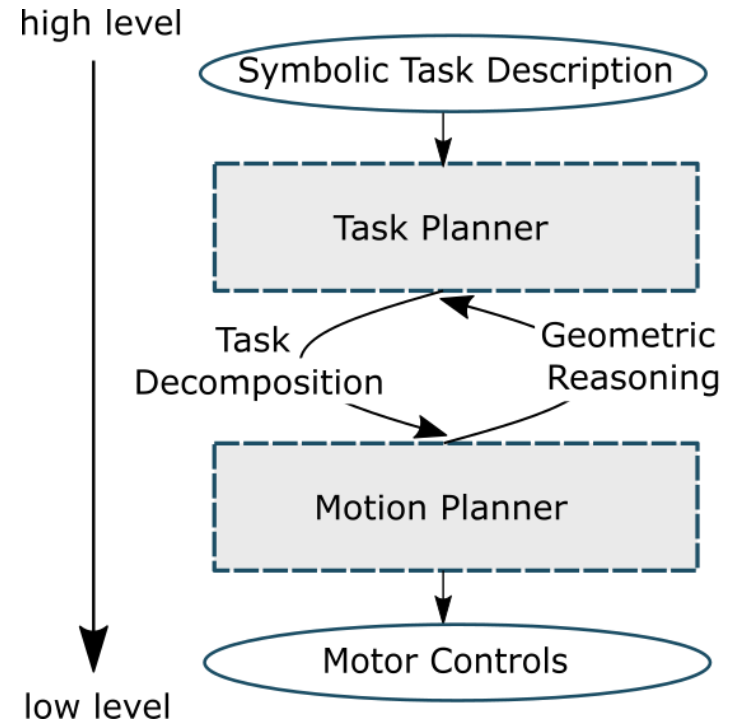
Why Task and Motion Planning?

Task Planner

- Efficient planning with logic descriptions for long time horizons

Motion Planner

- Considers the detailed geometric specification of the environment



Challenge:

Transform symbolic actions into feasible motion in variable scenarios



Related Work

→ Improve state-of-the art methods in terms of **efficiency**

TAMP framework	Limitation	Proposed solution	Efficiency gain
Search-based [Dantam, 2018] [Bidot, 2017]	Deliberation before every execution	Learning-based approach	Faster at execution time
Learning-based using RL [Quack, 2015]	Only performed in low-dimensional spaces	Apply RL on LfD	Efficient motion generation with LfD
Learning-based using LfD [Agostini, 2020]	Many demonstrations required	Generalization with RL	Only one demonstration required

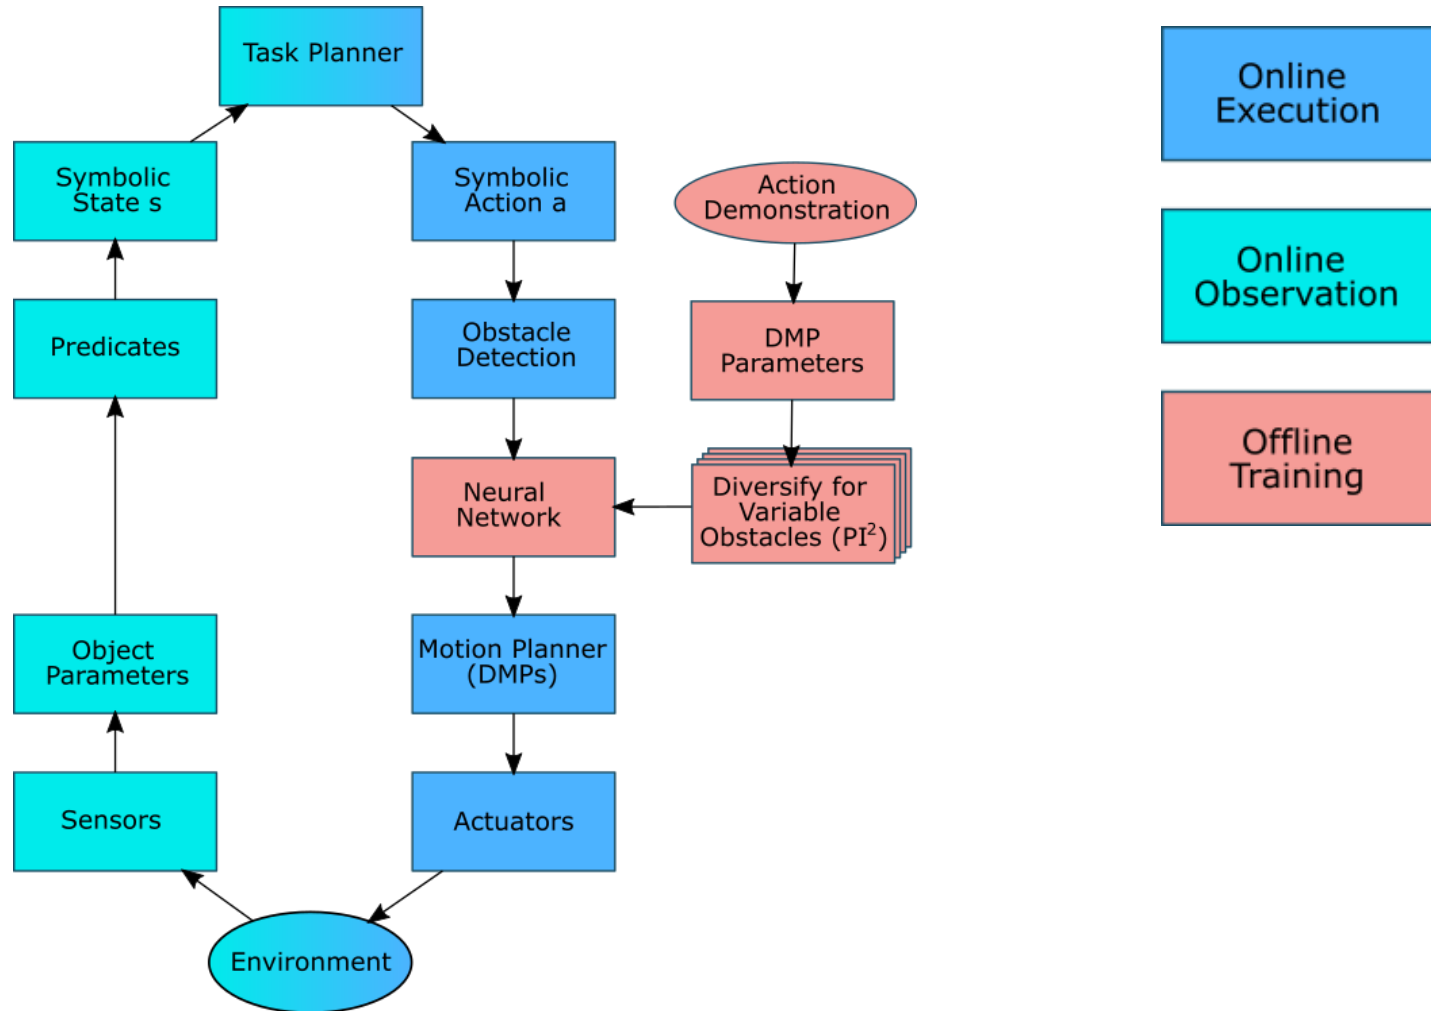


Contributions

- Our TAMP framework permits executing complex tasks comprising long action sequences with obstacle avoidance.
- Each symbolic action is grounded using DMPs that an action policy provides for variable object configurations.
- PI^2 efficiently generates diverse sets of optimal collision-free trajectories serving as training samples to learn the action policy encoded in a NN.



The Proposed TAMP Framework



Training Methods

Why DMPs?


- Efficient motion encoding from one demo
- Translation, dilatation, rotation invariance

Why PI²?


- Tunes parameterized policies like DMPs
- Numerically stable based on SOC

Why NN?


- Encode action policy for action selection



Offline Training – Encode a Demonstration




DMP parameters θ_i^j of the forcing term at $j = 1$



No ability to avoid obstacles

→ Apply PI^2 iteratively to tune the DMP parameters θ_i



Offline Training - PI² Optimization Step

The forcing term θ_i^j of the DMPs represents the current policy

1. Generate K random samples from the current policy θ_i^j


$$\theta_{i,k}^{j+1} = \theta_i^j + \epsilon_{i,k}^j$$

2. Weight each sample k using a cost function S

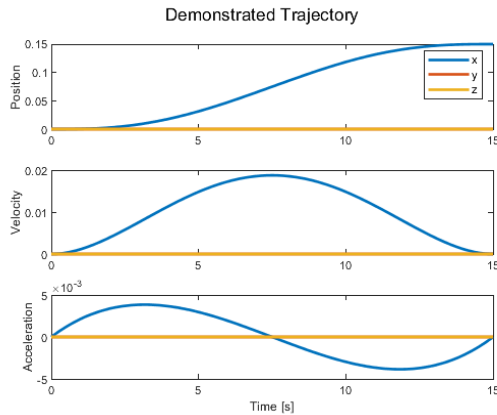
$$W_{\theta}(\theta_{i,k}^j) = \exp\left(-\gamma \frac{S(\theta_{i,k}^j) - \min S(\boldsymbol{\theta}_i^j)}{\max S(\boldsymbol{\theta}_i^j) - \min S(\boldsymbol{\theta}_i^j)}\right)$$

3. Compute the new policy θ_i^{j+1} as weighted average of the samples


$$\theta_i^{j+1} = \frac{\sum_{k=1}^K W_{\theta}(\theta_{i,k}^j) \theta_{i,k}^j}{\sum_{k=1}^K W_{\theta}(\theta_{i,k}^j)}$$




Offline Training – Diversify the Demonstration



DMP parameters θ_i^j of the forcing term at $j = 1$





$$S = -H = \min(h_{p1}, h_{p2})$$



4x

Tuned forcing term θ_i^j for $j = \{50, 100, \dots, J\}$



Defintion of the Trajectory Shape

Two variables specify the trajectory shape r_c, r_L


Degree of curvature $r_c = H/l$

$$r_c = H/l$$

Trajectory steepness $r_L = (|\tilde{L}(2) - \tilde{L}(1)| + 1)/B$

$$r_L = (|\tilde{L}(2) - \tilde{L}(1)| + 1)/B$$

→ r_c, r_L are dilatation invariant




Same set of forcing term parameters θ_i



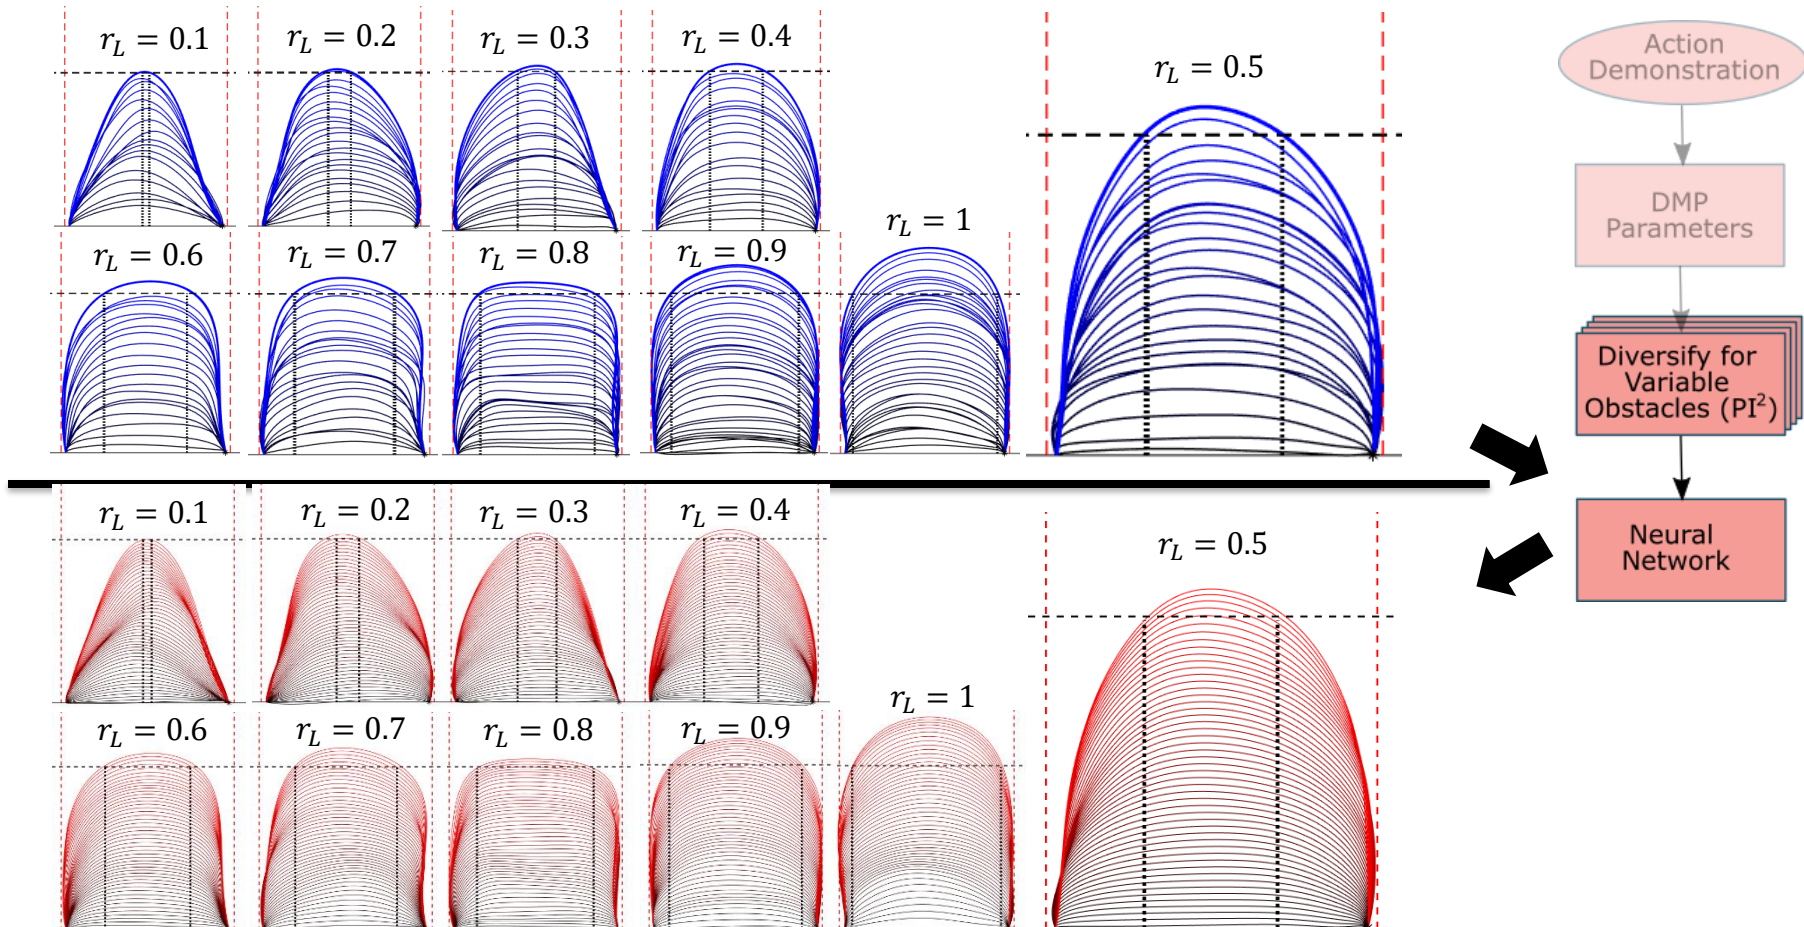
Same trajectory shape




Same pair of $r_C = 1, r_L = 0.4$




Offline Training – Learning Action Selection




→ Learned optimal parameters to ground symbolic actions in varying scenarios



Initial State



Goal State



- Cube
- Empty cell

Task Planner

Decompose the task into a sequence of *pickplace* actions

The *domain*

constants

air

predicates

(on ?cell ?cube)

symbolic action

(pickplace ?from ?to ?c)


precondition

(and (on ?from ?c)
(on ?to air))

effect

(and (on ?from air)
(on ?to ?c)
(not (on ?from c))
(not (on ?to air)))

Initial State



Task Plan

```
#1: pickplace cell11 cell15 cube4  
#2: pickplace cell13 cell11 cube8  
#3: pickplace cell19 cell13 cube7  
#4: pickplace cell11 cell19 cube8  
#5: pickplace cell18 cell11 cube2  
#6: pickplace cell16 cell18 cube6  
#7: pickplace cell17 cell16 cube1  
#8: pickplace cell15 cell17 cube4  
#9: pickplace cell14 cell15 cube3
```

The *task*

objects

cube1 cube2 ...
cube8
cell11 cell12 ...
cell19

init

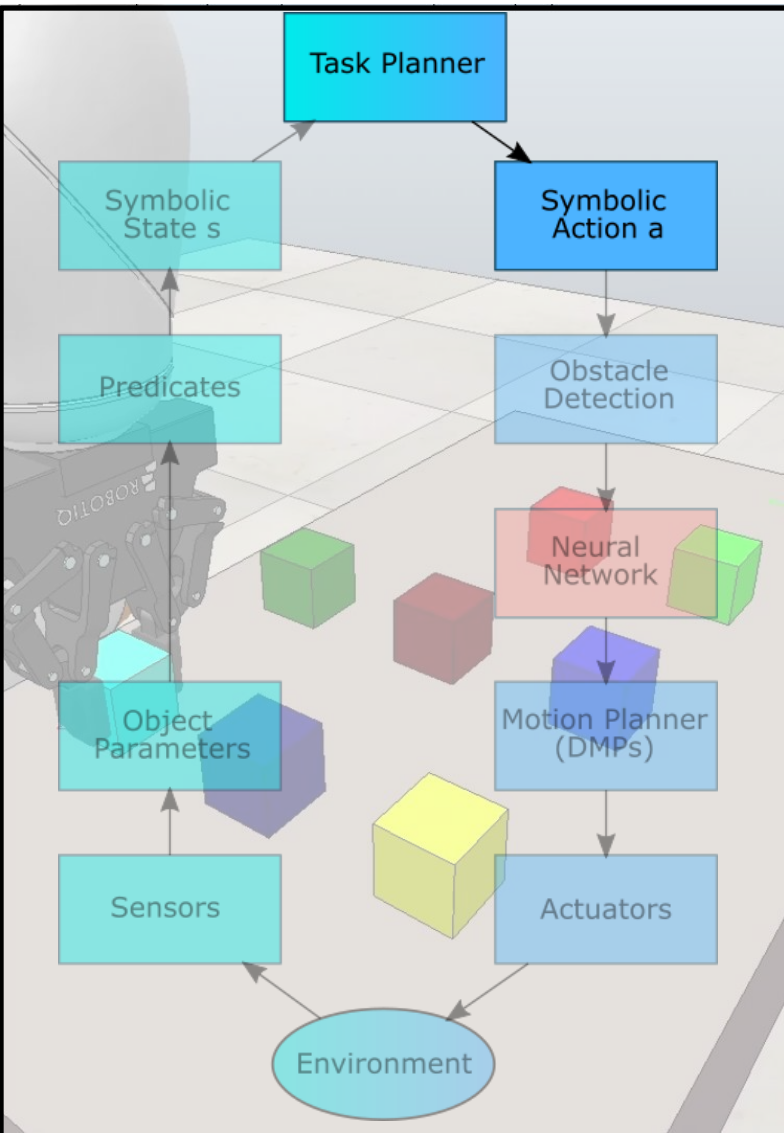
(on cell15 air)
(on cell11 cube4)
...

goal

(on cell14 air)
(on cell12 cube5)
...



Online Loop – Select Symbolic Action



Current symbolic action:
Pick and place *cube8*
from *cell3* to *cell1*

Task Plan


- #1: pickplace cell11 cell15 cube4
- #2: pickplace cell13 cell11 cube8
- #3: pickplace cell19 cell13 cube7
- #4: pickplace cell11 cell19 cube8
- #5: pickplace cell18 cell11 cube2
- #6: pickplace cell16 cell18 cube6
- #7: pickplace cell17 cell16 cube1
- #8: pickplace cell15 cell17 cube4
- #9: pickplace cell14 cell15 cube3

Pick

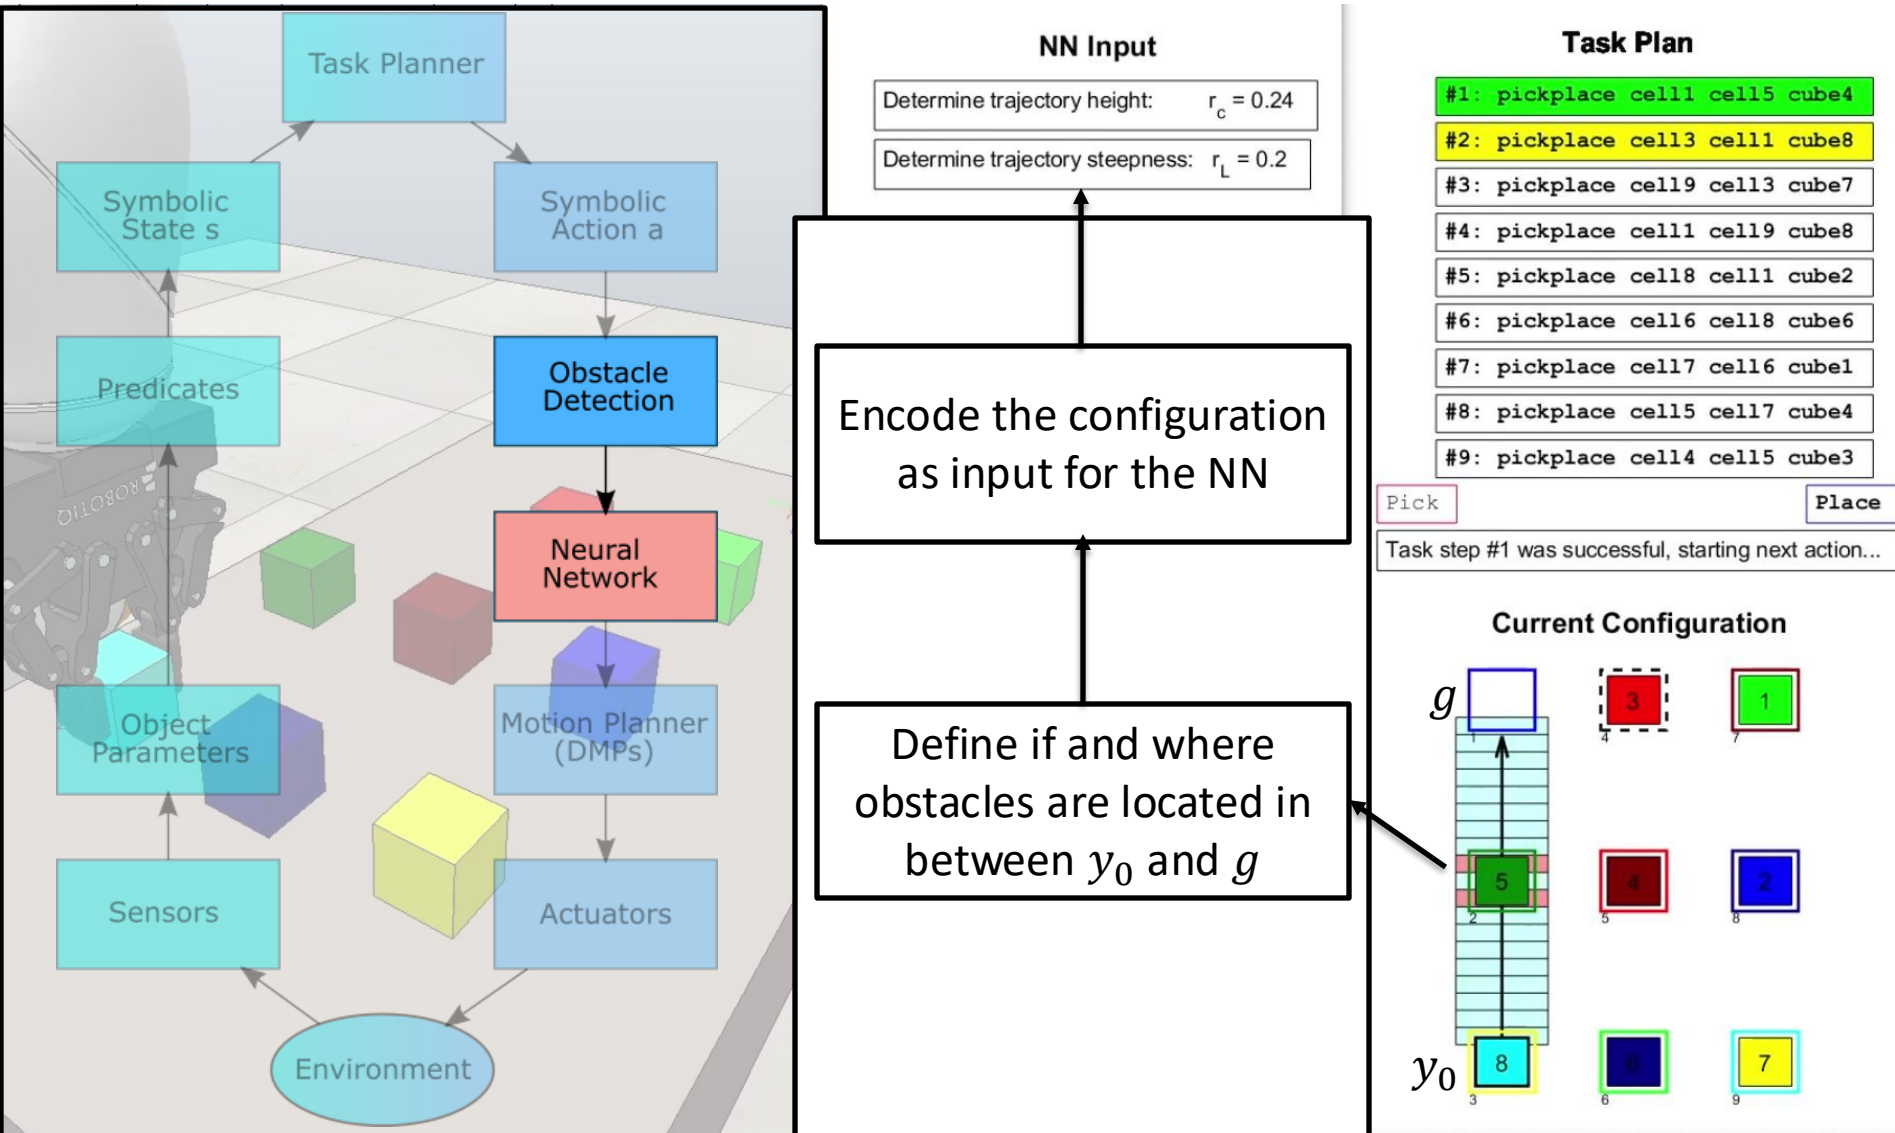
Place

Task step #1 was successful, starting next action...

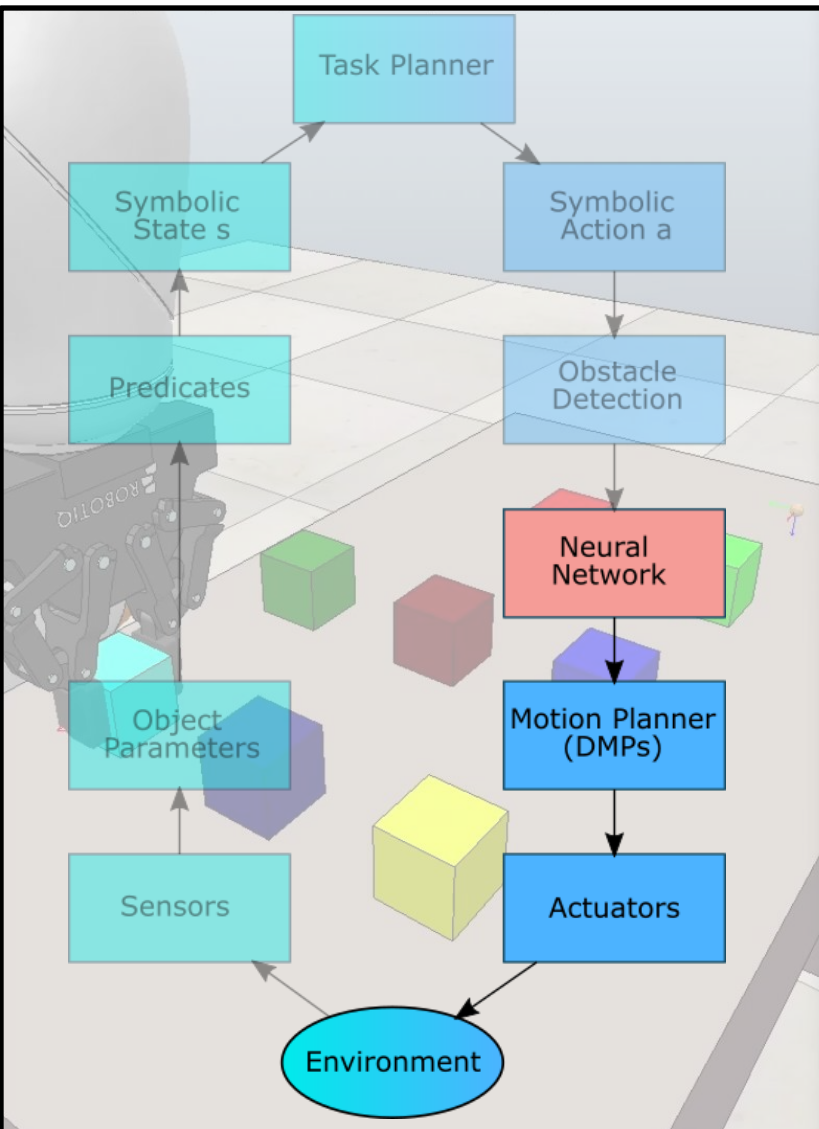
Current Configuration



Online Loop – Encode Current Configuration



Online Loop – Execute Action Policy





NN Input


Determine trajectory height: $r_c = 0.24$

Determine trajectory steepness: $r_L = 0.2$

NN Output



Generated Trajectory



Pick Place

Trajectory Precision

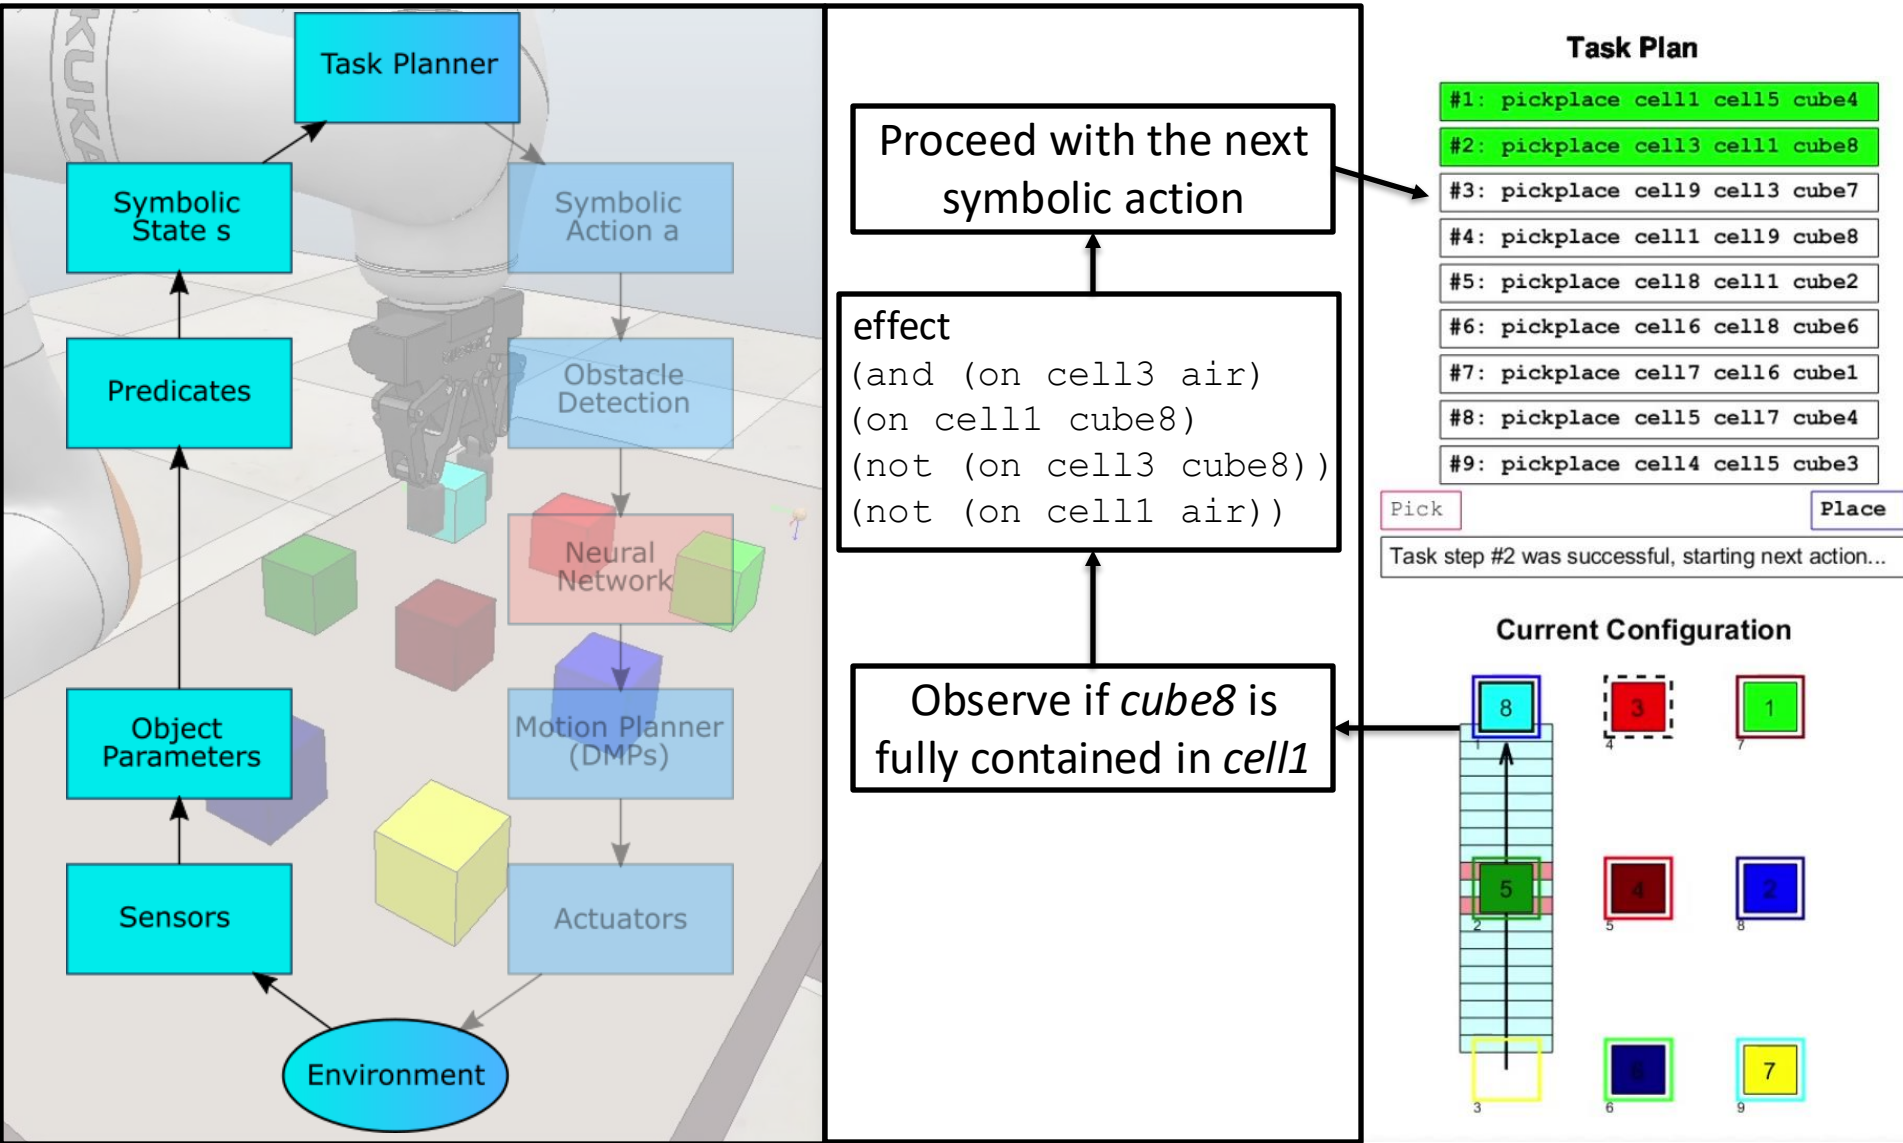
Goal deviation: $d_g = 0.02\%$

Height deviation: $d_H = -0.5\%$


Retrieve and set the forcing term parameters of the DMPs

Generate the trajectory and send it to the robot arm

Online Loop – Observe Changes to the Environment



Complete Execution of a Task Example




2x

NN Input

r_c
 r_L

NN Output



Generated Trajectory



Trajectory Precision

d_g
 d_H

Task Plan


- #1: pickplace cell11 cell15 cube4
- #2: pickplace cell13 cell11 cube8
- #3: pickplace cell19 cell13 cube7
- #4: pickplace cell11 cell19 cube8
- #5: pickplace cell18 cell11 cube2
- #6: pickplace cell16 cell18 cube6
- #7: pickplace cell17 cell16 cube1
- #8: pickplace cell15 cell17 cube4
- #9: pickplace cell14 cell15 cube3

Pick

Place

Starting the Task...

Current Configuration


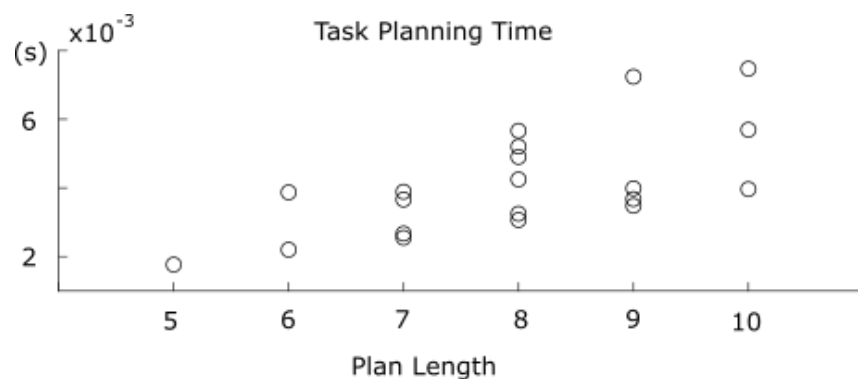


Experimental Evaluation


20 consecutive task executions with random initial and goal state

→ All 159 performed *pickplace* actions successful

Training Times	9 min
Generating the demonstration	0.02 s
PI ² optimizations	241 s
Training of the NN	279 s




Evaluation – Compared to [Dantam, 2018]



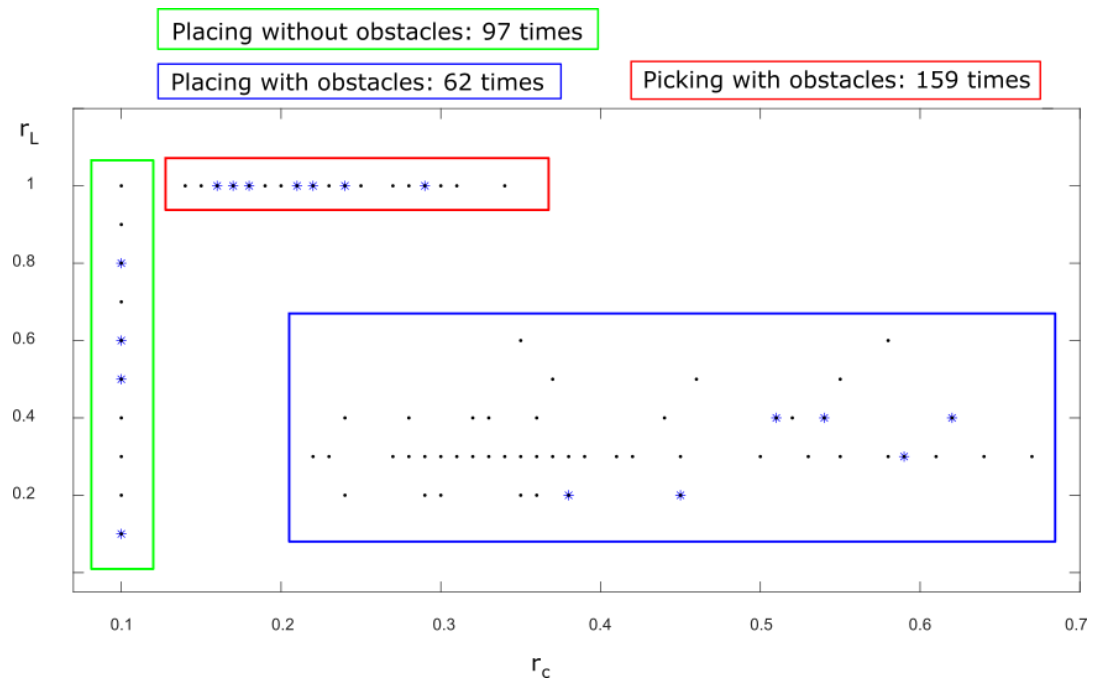
Generalization Ability

Varying block dimensions



Varying cell positions



- Learned action policy finds collision-free trajectories for all scenarios
- More trajectory shapes required



Complete Execution of a Task Example





NN Input

Determine trajectory height: $r_c = 0.26$

Determine trajectory steepness: $r_L = 1$

NN Output



Pick

Place

Trajectory Precision

Goal deviation: $d_g = 0.92\%$

Height deviation: $d_H = -7.39\%$

Task Plan

#1: pickplace cell11 cell19 cube3

#2: pickplace cell12 cell11 cube8

#3: pickplace cell14 cell12 cube1

#4: pickplace cell16 cell14 cube6

#5: pickplace cell15 cell16 cube5

#6: pickplace cell17 cell15 cube7

#7: pickplace cell11 cell17 cube8

#8: pickplace cell18 cell11 cube4

#9: pickplace cell13 cell18 cube2


#10: pickplace cell19 cell13 cube3

Pick

Place

Starting the Task...

Current Configuration



Limitation & Future Work

Limitations

1. Restricted to initial and goal positions in the same plane
2. Constant gripper orientation
3. Obstacle avoidance in one dimension only
4. Heuristics required for obstacle definition and optimization costs

Future Work

- Train NN on independent data to improve generalization
- Let RL agent learn to select appropriate input parameters and goal poses



Conclusion

- TAMP framework that utilizes LfD to efficiently generate motion and RL to generalize this motion
- PI^2 permits learning DMP parameters from a single demonstration to avoid obstacles of varying size and in varying situations
- After a few minutes of training, the action policy reliably selects collision-free trajectories to ground symbolic actions of a complex task



References



J. Bidot, L. Karlsson, F. Lagriffoul, A. Saffiotti

Geometric backtracking for combined task and motion planning in robotic systems.

In: *Artificial Intelligence*, 2017, pp. 229–265.



N. T. Dantam, Z. K. Kingston, S. Chaudhuri, L. E. Kavraki

An incremental constraint-based framework for task and motion planning.

In: *The International Journal of Robotics Research*, 2018, pp. 1134–1151



B. Quack, F. Wörgötter, A. Agostini

Simultaneous learning at different levels of abstraction.

In: *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)* , 2015, pp. 4600–4607.



A. Agostini, M. Saveriano, D. Lee, J. Piater

Manipulation planning using object-centered predicates and hierarchical decomposition of contextual actions.


In: *IEEE Robotics and Automation Letters*, 2020, pp. 5629–5636.



Appendix



Cost Function





$$S = -H + c_1 \cdot S_{prec} + c_2 \cdot S_{scope}$$

$$H = \min(h_{p1}, h_{p2})$$

$$S_{prec} = \|g - y_{end}\|$$

$$S_{scope} = - \sum_{t=1}^T \min(0, m + y_{X,t} - y_{0,X}) - \sum_{t=1}^T \min(0, m + g_X - y_{X,t})$$







NN Input

Determine trajectory height: $r_c = 0.24$

Determine trajectory steepness: $r_L = 1$

NN Output



Pick

Place

Trajectory Precision

Goal deviation: $d_g = 0.78\%$

Height deviation: $d_H = -6.29\%$

Task Plan

#1: pickplace cell11 cell15 cube4

#2: pickplace cell13 cell11 cube8

#3: pickplace cell19 cell13 cube7

#4: pickplace cell11 cell19 cube8

#5: pickplace cell18 cell11 cube2

#6: pickplace cell16 cell18 cube6

#7: pickplace cell17 cell16 cube1

#8: pickplace cell15 cell17 cube4



#9: pickplace cell14 cell15 cube3

Pick

Place





Starting the Task...

Current Configuration



DMP: Roto-Dilatation Invariance [Ginesi, 2019]

One demonstration encoded in three DMPs




New goal position (rotated around z: $\alpha=90^\circ$)

Same forcing term parameters

→ Unexpected trajectory shape




→ Goal position is not reached precisely



Rotate the forcing term parameters


$$f_x^{new} = \cos \alpha * f_x^{demo}$$

$$f_y^{new} = \sin \alpha * f_y^{demo}$$





$$\tilde{L} = (7,14)$$



$$\tilde{L} = (7,14)$$

